

Object Oriented Programming - Lecture 8

Diana Borza - diana.borza@ubbcluj.ro

April 9, 2024

- Move semantics
- RAII
- Smart pointers
- threads

L-values and R-values I

Each C++ expression (an operator with its operands, a literal, a variable name, etc.) is characterized by two independent properties: *a type* and *a value category*. In C++ there are several types of *value categories*:

- ❶ l-values - so-called, historically, because lvalues could appear on the left-hand side of an assignment expression;
- ❷ r-values - so-called, historically, because rvalues could appear on the right-hand side of an assignment expression;
- ❸ gl-value - “generalized” lvalue;
- ❹ pr-value - “pure” rvalue;
- ❺ xvalue - an “eXpiring” value

In this course, we'll just cover l-values and r-values. For further information, you can check:

https://en.cppreference.com/w/cpp/language/value_category

L-values and R-values II

- l-values

- "values that are suitable to be on the left-hand side of an assignment expression";
- all l-values have assigned memory addresses;
- a function or an object (or an expression that evaluates to a function or object).

- r-values

- "everything that is not an l-value";
- e.g. literals (10, 'a' etc.), temporary values ($x+3$), anonymous objects (Student" John Smith")
- r-values are typically evaluated for their values
- r-values have expression scope (they die at the end of the expression they are in)
- typically r-values cannot be assigned to.

Select the l-values and the r-values from the following code snippet:

```
int a = 2;  
int b = 3;  
int sum = a + b;  
MyClass cl = MyClass(5, 7);  
int d = myFunction();
```

```
int a = 2; // a is an lvalue, 2 is an rvalue
int b = 3; // b is an lvalue, 3 is an rvalue
int sum = a + b; // sum is an l value, a+b is an rvalue
MyClass cl = MyClass(5, 7); // cl is an lvalue, MyClass(5, 7) is an rvalue (anonymous object)
int d = myFunction(); // d is an lvalue
           // the temporary value returned by myFunction()
           // (a function which returns an int ) is an rvalue
```

- Since C++11 it is possible to have a reference for an r-value.
- An r-value reference is a reference that is designed to be initialized with an r-value (**only**).
- While an l-value reference is created using a single ampersand, an r-value reference is created using a double ampersand:
 - `int x{5}; int &xRef{x}; // xRef is a l-value reference`
 - `int &&rRef {5}; // rref is a r-value reference`
- r-value references **extend the lifespan** of the object they are initialized with to the lifespan of the r-value reference;
- (non-const) r-value references allow you to modify the r-value.

L-values references and R-values references

- L-value references

L-value reference	Can be initialized with	Can modify
Modifiable l-values	Yes	Yes
Non-modifiable l-values	No	No
R-values	No	No

- L-value references to const

L-value reference to const	Can be initialized with	Can modify
Modifiable l-values	Yes	No
Non-modifiable l-values	Yes	No
R-values	Yes	No

L-values references and R-values references

- R-value references

R-value reference	Can be initialized with	Can modify
Modifiable l-values	No	No
Non-modifiable l-values	No	No
R-values	Yes	Yes

- R-value references to const

R-value reference to const	Can be initialized with	Can modify
Modifiable l-values	No	No
Non-modifiable l-values	No	No
R-values	Yes	No

- If we construct an object or do an assignment where the argument is an **l-value**, the only thing we can reasonably do is **copy** the l-value.
 - We cannot alter the l-value as it may be used again in the program.
- If we construct an object or do an assignment where the argument is an **r-value**, then we know that r-value is just a temporary object of some kind.
 - Instead of copying it (which can be expensive), we can simply **transfer its resources** (which is cheap) to the object we're constructing or assigning.
 - the temporary will be destroyed at the end of the expression anyway, so we know it will never be used again.

Move constructor and move assignment

- The move constructor and move assignment are called when those functions have been defined, and **the argument for construction or assignment is an r-value** (literal or temporary value).
- A move constructor and move assignment operator will NOT be provided by default.

Move constructor

- Whereas the goal of the copy constructor and assignment op. is to make a copy of one object to another, the goal of the move constructor and assignment op. is to **move ownership** of the resources from one object to another (which is much less expensive than making a copy).
- syntax: `class_name(class_name &&);`
- Disabling copy constructor: use **delete** keyword:
 - If, instead of a function body, the special syntax `=delete ;` is used, the function is defined as *deleted*.
 - Any use of a deleted function is ill-formed (the program will not compile).

Move assignment

- Move assignment operators typically "**steal**" the resources held by the argument (e.g. pointers to dynamically-allocated objects, file descriptors, TCP sockets, I/O streams, running threads, etc.), rather than make copies of them;
- they leave the argument in some valid but otherwise indeterminate state.
- syntax:
- syntax: `class_name & class_name :: operator=(class_name &&);`

Demo

Move constructor and move assignment.

Rule of five I

- The rule of five is a modern expansion of the rule of three.
- Remember the rule of three?

Rule of five II

- The rule of three specifies that if a class implements any of the following functions, it should implement all of them:
 - copy constructor;
 - copy assignment operator;
 - destructor.
- **In addition**, the rule of five identifies that it is usually appropriate to also provide the following functions to allow for optimized copies from temporary objects:
 - move constructor;
 - move assignment operator.

- `std::move` is a library function to convert its argument into an r-value.
- We pass an l-value to `std::move`, and it will return an r-value reference.
- `std::move` is defined in the `<utility>` header.

- RAII - *Resource Acquisition Is Initialization* - is a programming technique in which a resource is tied to the lifetime of objects which acquires it.
- A resource can be memory location, a database connection, a file, a network socket etc.
- The resource is *acquired* (allocated) during object creation (specifically *initialization*), by the constructor, while resource *deallocation* (release) is done during object destruction (specifically *finalization*), by the destructor.
 - *Constructor Acquires, Destructor Releases (CADRe)*
- RAII is used to **avoid resource leaks** and to write **exception-safe code**.

- Advantages over garbage collection (from other programming languages):
 - RAII is a programming idiom which offers automatic management for different kinds of resources, not just memory.
 - The runtime environment is faster, as there is no separate mechanism involved (like the garbage collector).

- You already used RAII in your programs:
 - 1 Files in C++:
 - When using an object of type ifstream or ofstream, the constructor will automatically open the file.
 - When the object gets destroyed, the destructor automatically closes the file.
 - 2 The STL containers manage memory using the RAII programming technique.

Implementing RAI - I

- The lifetime of the objects allocated on the stack is automatically managed by the compiler
 - The compiler automatically invokes constructors to initialize objects;
 - The compiler automatically calls the destructor when the object goes out of scope.
- ❶ Create a wrapper for your resource, based on the following rule: allocate the resource in constructor, release it in the destructor.
- ❷ Use the resource (directly) wherever you need it.
- ❸ The resource will be automatically deallocated when the wrapper's scope is left.

Demo

RAII - smart pointer


Smart pointers in STL

- A smart pointer is a composition class that is designed to manage dynamically allocated memory, and ensure that memory gets deleted when the smart pointer object goes out of scope.
- Smart pointers should never be dynamically allocated themselves. Why?
- Smart pointers are defined in the `<memory>` header.
- In STL there are 3 types of smart pointers:
 - `std::unique_ptr`;
 - `std::shared_ptr`;
 - `std::weak_ptr`.

- It is used to manage any dynamically allocated object that is not shared by multiple objects.
- It retains *exclusive ownership* of the object, it does not share the object.
- A std::unique_ptr **cannot be**:
 - copied to another unique_ptr;
 - passed by value to a function;
 - used in any STL algorithm that requires copies to be made.
- A std::unique_ptr **can only** be **moved**: `std::move()`.

std::unique_ptr II

```
auto ptrA = make_unique<Song>(L"Diana Krall", L"The Look of Love");
```



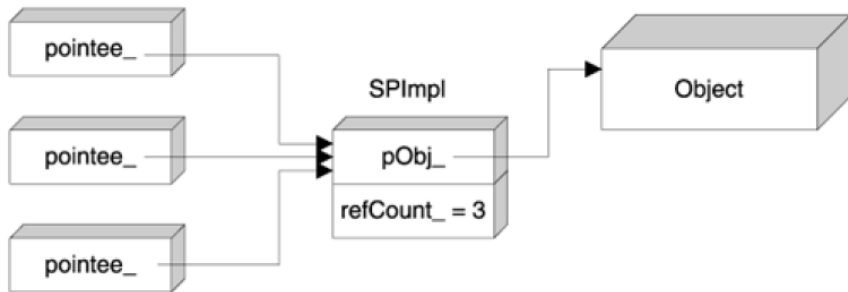
```
auto ptrB = std::move(ptrA);
```



std::unique_ptr III

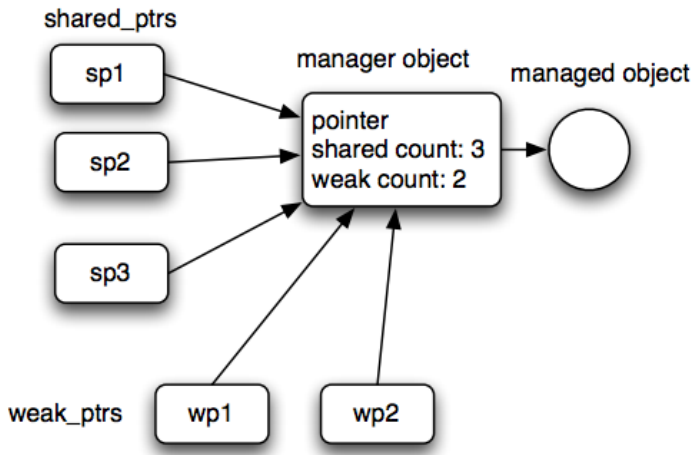
- `std::unique_ptr` provides two overloaded operators to access the raw pointer: `operator*` (reference to the resource) and `operator— >` (pointer to the reference).
- checking that `std::unique_ptr` actually holds a resource: `std::unique_ptr` has a cast to `bool` that returns true if the `std::unique_ptr` is managing a resource.
- It is recommended that `std::unique_ptr` are created via the `std::make_unique()` function.
- To pass `std::unique_ptr` to a function you should use `std::move()` → !!! the function will take ownership of the object. Alternative: pass the resource itself to the function.

- std::shared_ptr was developed for the cases where you need multiple smart pointers *co-owning* a resource.
- Internally, std::shared_ptr uses **reference counting**. Reference counting keeps track of the **number of smart pointers** pointing to (or owning) the same object.
- If at least one std::shared_ptr is pointing to the resource (i.e. the reference number is greater than 0), the resource will not be deallocated.
- When the reference number becomes zero, the pointee object is deleted.



- std::shared_ptr can be copied and moved (move transfers ownership).
- has more overhead than unique_ptr (because of the internal reference counting), therefore, whenever possible, prefer unique_ptr.
- it should be constructed with the std::make_shared function.
- can be created from a std::unique_ptr, but not the otherway around.
- Problems with shared pointers: **cyclic dependencies**.
- A circular reference (cyclical reference, cyclical dependency) refers to a series of references where each object references the next, and the last object references back to the first, causing a *referential loop*.

- std::weak_ptr was designed to solve the cyclical dependency problem which can occur when using shared pointer.
- Can be seen like an observer – it can observe and access the same object as a std::shared_ptr, but **it is not considered an owner**.
- Used to access the underlying object of a std::shared_ptr without causing the reference count to be incremented.



- You are using two classes to capture the information about a soccer team. You have 2 classes *Team* and *Member*.
 - A team has pointers to its members.
 - Each member can have a pointer to the team it belongs to.
 - If all pointers (to members and to team) are std::shared_ptr, what happens when the team goes out of scope? → memory leak.
 - Therefore, the members should have a weak pointer to their team.

Smart pointers - best practices

"Smart developers use smart pointers." - Kate Gregory

"In modern C++, raw pointers are only used in small code blocks of limited scope, loops, or helper functions where performance is critical and there is no chance of confusion about ownership." -

<https://docs.microsoft.com/en-us/cpp/cpp/smart-pointers-modern-cpp?redirectedfrom=MSDN&view=vs-2019>

Smart pointers - advantages

- Smart pointers increase productivity and improve the robustness of the program.
- The programmer does not need to be concerned with memory management (provided the smart pointers are used correctly).
- Avoid memory leaks.
- Write exception safe code.

Threads



- **Program** is the code that is stored on your computer that is intended to fulfill a certain task; stored on disk or in non-volatile memory in a form that can be executed by your computer (machine language).
- A **process** is a program that has been loaded into memory along with all the resources it needs to operate.
- Each process has a separate memory address space (a process runs independently and is isolated from other processes).
- Switching from one process to another is relatively slow: saving and loading registers, memory maps, and other resources.

Resources of a process

- “registers” - data storage locations which are part of the computer processor (CPU).
- ”program counter” / “instruction pointer” - keeps track of the current location of the program sequence.
- stack? heap?




Figure source:

<https://www.backblaze.com/blog/whats-the-diff-programs-processes-and-threads/>

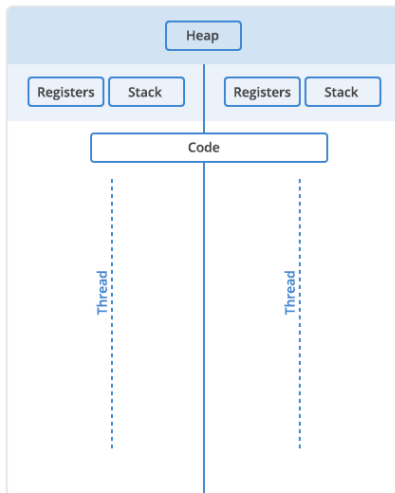
- A thread is the unit of execution within a process.
- A process can have anywhere from just one thread (single-threaded app.) to many threads (multi-threaded app.).
- When a process starts, it is assigned memory zone and resources.
- Each thread in the process shares that memory and resources.
- Each **thread will have its own stack**, but all the threads in a process will **share the heap**.
- The disadvantage is that a problem with one thread in a process will certainly affect other threads and the viability of the process itself.

Threads vs. Processes I

Single Thread



Multi Threaded



Threads vs. Processes II

Threads

- Threads use the memory of the process they belong to;
- Inter-thread communication can be faster than IPC (threads share memory with the process they belong to);
- Context switching between threads of the same process is less expensive;
- Threads are lighter weight operations.

Processes

- Each process has its own memory space;
- Inter-process communication is slow as processes have different memory addresses;
- Context switching between processes is more expensive;
- Processes don't share memory with other processes;
- Processes are heavyweight operations.

Parallelism and concurrency I


- On a system with multiple processors or CPU cores, multiple processes or threads can be executed in parallel.
- On a single processor, though, it is not possible to have processes or threads truly executing at the same time.
 - the CPU is shared among running processes or threads using a process scheduling algorithm that divides the CPU's time and yields the **illusion** of parallel execution;
 - the time given to each task is called a *time slice*.
- **Parallelism** - genuine simultaneous execution.
- **Concurrency** - interleaving of processes in time to give the appearance of simultaneous execution.

Parallelism and concurrency II

Concurrency



Parallelism




Threads in C++

- `std::thread` - class to represent individual threads of execution.
- declared in the header `<thread>` .
- The constructor of `std::thread` can take as parameters:
 - a pointer to function that will be executed by the thread;
 - arguments passed to the call of this function (optional).
- The `join()` function from the `std::thread` class returns when the thread execution has completed.
 - This blocks the execution of **the thread that calls this function** until the function called on construction returns (if it hasn't yet).

Thread synchronization I

- Synchronization - refers to one of the following concepts:
 - **Process synchronization** refers to the idea that multiple processes are to join up or handshake at a certain point, in order to reach an agreement or commit to a certain sequence of action.
 - **Data synchronization** refers to the idea of keeping multiple copies of a dataset in coherence with one another, or to maintain data integrity.

Multithreaded programming



Thread synchronization III

- **Mutual exclusion** algorithms prevent multiple threads from simultaneously accessing shared resources. This prevents data races and provides support for synchronization between threads.
- <https://en.cppreference.com/w/cpp/thread/mutex>
- https://en.cppreference.com/w/cpp/thread/lock_guard - a mutex wrapper that provides a convenient RAII-style mechanism for owning a mutex for the duration of a scoped block.




Mutex idea

`door.lock()`



Good morning **Mr. ThreadLibrary!**
Can I visit Vector V? I need to give him something.

Yes, no one is in there. Remember to **unlock the door** and leave as soon as you can! Other people want to see him.




after some nanoseconds . . .

`door.lock()`



Good morning **Mr. ThreadLibrary!**
Can I see Vector V?

Sorry, the **door is locked**. Wait here until the other guy leaves.



Other synchronization mechanisms

- Other synchronization mechanisms (not a subject of this lecture):
 - semaphores: https://en.cppreference.com/w/cpp/thread/counting_semaphore
 - condition variables: https://en.cppreference.com/w/cpp/thread/condition_variable.